Qualifying Paper Report for Data Fission: Splitting A Single Data Point

Naitong Chen

January 20, 2023

1 Problem Definition

Given a dataset $(X_i)_{i=1}^n \overset{\text{i.i.d.}}{\sim} \pi_\theta$, with π_θ being a distribution from the exponential family whose parameter θ is of interest. We decompose each X_i to $f(X_i)$ and $g(X_i)$ such that both parts contain information about θ, and there exists some function h such that $X_i = h(f(X_i), g(X_i))$ satisfying either of the two properties:

- (P1): $f(X_i)$ and $g(X_i)$ are independent with known distributions (up to unknown θ);
- (P2): $f(X_i)$ has a known marginal distribution and $g(X_i)$ has a known conditional distribution given $f(X_i)$ (up to unknown θ).
  Typos and missing terms in the proof of the original paper. See Theorem A.2 for a modified version.

2 Significance

3 Limitations and challenges

- paper does not discuss robustness of the method to the distribution assumptions
- experiments do not cover cases where fissioned data get transformed to follow a different distribution
- paper does not discuss how to choose tuning parameter (controlling amount of information split between $f(X)$ and $g(X)$)
- following the previous point, this paper does not discuss the relations between having a discrete vs. continuous tuning parameter (e.g. the two different ways of fissioning exponentially distributed data in Appendix B of Leiner et al. (2022))
- Proofs have missing terms, inaccurate notations.
- The discrete version of data fission for Gamma data may be incorrect.

4 Paper-specific project

I noticed that most of the experiments and simulation studies in the paper only cover cases where the distributions of $f(X)$ and $g(X) \mid f(X)$ are in the same family as the original data X (i.e. Gaussians or Poissons with different parameters), and there is minimal discussion on when this is not the case. I would like to therefore focus my paper-specific project on an instance of data fission where $f(X)$ follows a distribution that is not in the same family as X or $g(X) \mid f(X)$.

The particular case that I would like to focus on is to **construct selective CIs in the fixed-design GLM model** where $Y_i \overset{\text{i.i.d.}}{\sim} \mathsf{Gam}(\alpha, \exp(\theta_i))$ where $\theta_i = \beta^\top x_i$ (listed under Appendix B of Leiner et al. (2022)).

- For each Y_i with $i \in \{1, \dots, n\}$, draw $Z_i = (Z_{i1}, \dots, Z_{iB})$ where each element is i.i.d. $Z_{ib} \sim \mathsf{Poiss}(Y_i)$ with $b \in \{1, \cdots, B\}$.
  Then $f(Y_i) = Z_i$, where each element is i.i.d. $\mathsf{NB}\left(\alpha, \frac{\exp(\theta_i)}{\exp(\theta_i)+B}\right)$,
  $g(Y_i) \mid f(Y_i)$ has conditional distribution $\mathsf{Gam}(\alpha + \sum_{b=1}^{B}[f(Y_i)]_b, \exp(\theta_i) + B)$.
  I believe the marginal distribution of $f(Y_i)$ is incorrect. See Example A.3.
- For each Y_i, draw $Z_i \sim \mathsf{Poiss}(\tau Y_i)$ with $\tau \in (0, \infty)$.
  Then $f(Y_i) = Z$, where each element is i.i.d. $\mathsf{NB}\left(\alpha, \frac{\theta_i}{\theta_i+\tau}\right)$,
  $g(Y_i) \mid f(Y_i)$ has conditional distribution $\mathsf{Gam}(\alpha + f(Y_i), \theta_i + \tau)$.

The proposed procesure is

- Decompose each y_i using one of the two above procedures (assuming we are going with the second one)
- fit $f(y_i)$ by maximizing

$$\sum_{i=1}^{n} \log \mathsf{NB}\left(z_i \mid \alpha, \frac{\beta^\top x_i}{\beta^\top x_i + \tau}\right) + \lambda\|\beta\|_1$$

  (I've verified that this function is convex in β.)
- Fit a Gamma GLM model using just the selected features from the previous step
- Since Gam is in the exponential dispersion family, we can follow the exact same setup as in Appendix A.5 of Leiner et al. (2022) to use the QMLE procedure to construct CIs in the inference step.
  This function may be convex but without an argmin, or non-convex, depending on the values of y, z, and α. See Examples A.5 and A.6. This is already an undesirable quality of data fission.

Key areas that I would like to explore (through simulations)

- compare the variables selected using decomposed data following geometric distribution against those using data splitting and the (invalid) approach of using the same dataset twice for both selection and inference. The setup that I have is mostly inline with the simulation studies in Sections 4 and 5 in Leiner et al. (2022), however, I would like to ensure there is no influential point and all assumptions are met. This way we can better isolate the effect of transforming the original dataset to something that follows a different distribution.
- compare the two data fission procedures laid out above with discrete and continous tuning parameters (B and τ) for deciding how much information is split between $f(Y)$ and $g(Y) \mid f(Y)$. The second version seems like a continuous relaxation (under expectation) of the first data fission method. However, in the case with a discrete tuning parameter, the dimension of $f(Y_i)$ changes and we need to somehow account for that in the selection step. This is not addressed directly in the paper, but I think a natural way to deal with this is to think of stacking the elements of $f(Y_i)$ so that for each particular set of covariates x_i, we have multiple corresponding responses instead of 1. I would like to explore the connection between

these two fission processes (probably empirically) in terms of the amount of information allocated in each component of the fissioned data. The same simulation setup from the previous bullet point can be used here.
I believe this is no longer necessary because the result generalizing to $B > 1$ does not hold?

Revised plan
We investigate to what extent would we like the distributions of $X, f(X), g(X) \mid f(X)$ be similar to each other, in terms of parametric family and whether the parameters are random. For all of these comparisons, pick the tuning parameter so information allocated to $f(X)$ is the same compared to the counterpart in data splitting (ideally half).

- compare the first two versions of data fission for Gaussian linear regression against data splitting. All three parts follow Gaussian distributions, but one version has the value of Z being a parameter for the distribution of $g(X) \mid f(X)$, while the other does not. This can help us study the effect of having the parameter of the distribution of $g(X) \mid f(X)$ depending on Z. My hypothesis is that there will be higher variability on the width of the CIs.

- compare the two versions of data fission for Possion regression against data splitting. One version has all of $X, f(X), g(X) \mid f(X)$ following Possion distributions, the other version has $g(X) \mid f(X)$ following a binomial distribution and the others following Possion distributions. With this we can check the effect of not having the same parametric family of distributions. Use large enough samples with we are closer to asymptotic normality. But this effect may be confounded by the fact that the Binomial version has Z in the parameter while the other does not.

- both of the previous examples check the effect on the inference step. The Gamma example on the previous page has $f(X)$ following a different distribution. We can use this to study the effect of disimilarity between $f(X)$ and X on variable selection. But the corresponding MLE for the subsequent inference step is hard to optimize. If we use a working model, we can also then compare the inferential results against data splitting.

5 Discussion

References

James Leiner, Boyan Duan, Larry Wasserman, and Aaditya Ramdas. Data fission: splitting a single data point. *arXiv preprint arXiv:2112.11079 v4*, 2022.

A Proofs and examples

Example A.1. *Data fission can be viewed as a continuous analog of data splitting in terms of the allocation of Fisher information.*
Let $\{X_i\}_{i=1}^n$ *be i.i.d.* $\mathcal{N}(\theta, \sigma^2)$*. Let* $X := \left[X_1, \ldots, X_n\right]^\top$*. Recall that data splitting defines* $f(X)$ *and* $g(X)$ *as*

$$f^{split}(X) = \left[X_1, \ldots, X_{an}\right], \quad g^{split}(X) = \left[X_{an+1}, \ldots, X_n\right],$$

for $a \in \left\{\frac{1}{n}, \frac{2}{n}, \ldots, 1\right\}$*. Note that*

$$\mathcal{I}_{f^{split}(X)} = an\frac{1}{\sigma^2}, \quad \mathcal{I}_{g^{split}(X)} = (1-a)n\frac{1}{\sigma^2}.$$

On the other hand, data fission first simulates $\{Z_i\}_{i=1}^n$ *distributed as i.i.d.* $\mathcal{N}(0, \sigma^2)$ *and have, for some fixed* $\tau \in (0, \infty)$*,*

$$f^{fisson}(X) = \left[X_1 + \tau Z_1, \ldots, X_n + \tau Z_n\right], \quad g^{fisson}(X) = \left[X_1 - \tfrac{1}{\tau} Z_1, \ldots, X_n - \tfrac{1}{\tau} Z_n\right].$$

Note that for all $i \in \{1, \ldots, n\}$*,* $X_i + \tau Z_i \sim \mathcal{N}\left(\theta, (1+\tau^2)\sigma^2\right)$*,* $X_i - \frac{1}{\tau} Z_i \sim \mathcal{N}\left(\theta, (1+\frac{1}{\tau^2})\sigma^2\right)$*. We then have*

$$\mathcal{I}_{f^{fisson}(X)} = n\frac{1}{(1+\tau^2)\sigma^2}, \quad \mathcal{I}_{g^{fisson}(X)} = n\frac{1}{(1+\frac{1}{\tau^2})\sigma^2}.$$

By setting $a = \frac{1}{1+\tau^2}$*, we have* $\mathcal{I}_{f^{split}(X)} = \mathcal{I}_{f^{fisson}(X)}$ *and* $\mathcal{I}_{g^{split}(X)} = \mathcal{I}_{g^{fisson}(X)}$*.* ◁

Theorem A.2. *Suppose that for some* $A(\cdot), \phi(\cdot), m(\cdot), \theta_1, \theta_2, H(\cdot, \cdot)$*, the density of* X *is given by*

$$p(x \mid \theta_1, \theta_2) = m(x)H(\theta_1, \theta_2)\exp\{\theta_1^\top \phi(x) - \theta_2^\top A(\phi(x))\}.$$

Suppose also that there exists $h(\cdot), T(\cdot), \theta_3$ *such that*

$$p(z \mid x, \theta_3) = h(z)\exp\{\phi(x)^\top T(z) - \theta_3^\top A(\phi(x))\}$$

is a well-defined distribution. First, draw $Z \sim p(z \mid X, \theta_3)$*, and let* $f(X) := Z$ *and* $g(X) := X$*. Then,* $(f(X), g(X))$ *satisfy the data fission property (P2). Specifically, note that* $f(X)$ *has a known marginal distribution*

$$p(z \mid \theta_1, \theta_2, \theta_3) = h(z)\frac{H(\theta_1, \theta_2)}{H(\theta_1 + T(z), \theta_2 + \theta_3)},$$

while $g(X)$ *has a known conditional distribution given* $f(X)$*, which is*

$$p(x \mid z, \theta_1, \theta_2, \theta_3) = p(x \mid \theta_1 + T(z), \theta_2 + \theta_3).$$

Proof. Note that because the density $p(z \mid x, \theta_3)$ must integrate to 1, we can view the function $H(\theta_1, \theta_2)$ as a normalization factor since

$$H(\theta_1, \theta_2) = \frac{1}{\int_{-\infty}^{\infty} m(x)\exp\{\theta_1^\top \phi(x) - \theta_2^\top A(\phi(x))\}dx}.$$

Therefore, to compute the marginal density, we have

$$\begin{aligned} p(z \mid \theta_1, \theta_2, \theta_3) &= \int_{-\infty}^{\infty} m(x)h(z)H(\theta_1, \theta_2)\exp\{(T(z) + \theta_1)^\top \phi(x) - (\theta_2 + \theta_3)^\top A(\phi(x))\}dx \\ &= h(z)\frac{H(\theta_1, \theta_2)}{H(\theta_1 + T(z), \theta_2 + \theta_3)}. \end{aligned}$$

Similarly, the computation of the conditional density is straightforward

$$\begin{aligned} p(x \mid z, \theta_1, \theta_2, \theta_3) &= \frac{m(x)h(z)H(\theta_1, \theta_2)\exp\{(T(z) + \theta_1)^\top \phi(x) - (\theta_2 + \theta_3)^\top A(\phi(x))\}}{h(z)\frac{H(\theta_1, \theta_2)}{H(\theta_1 + T(z), \theta_2 + \theta_3)}} \\ &= m(x)H(\theta_1 + T(z), \theta_2 + \theta_3)\exp\{\phi(x)^\top (\theta_1 + T(z)) - (\theta_2 + \theta_3)^\top A(\phi(x))\} \\ &= p(x \mid \theta_1 + T(z), \theta_2 + \theta_3). \end{aligned}$$

This completes the proof. □

Example A.3. *Suppose* $X \sim \mathsf{Gam}(\alpha, \beta)$. *Draw* $Z = (Z_1, \cdots, Z_B)$ *where each element is i.i.d.* $Z_i \sim \mathsf{Poiss}(X)$ *and* $B \in \{1, 2, \dots\}$ *is a tuning parameter. Let* $f(X) = Z$ *and* $g(X) = X$.
By writing

$$\mathsf{Gam}(x \mid \alpha, \beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \exp\{(\alpha - 1)\log x - \beta x\},$$

where $\mathsf{Gam}(\cdot \mid \alpha, \beta)$ *denotes the pdf of* $\mathsf{Gam}(\alpha, \beta)$, *we have that* $\theta_2 = \beta, \theta_1 = \alpha - 1, \phi(x) = \log x, A(\phi(x)) = \exp(\phi(x)) = \exp(\log x) = x, m(x) = 1$. *Therefore,* $H(\theta_1, \theta_2) = \frac{\theta_2^{(\theta_1+1)}}{\Gamma(\theta_1+1)}$. *Now since*

$$\begin{aligned}
\mathsf{Poiss}(z \mid x) &= \prod_{i=1}^{B} \frac{1}{z_i!} \exp\{z_i \log x - x\} \\
&= \left(\prod_{i=1}^{B} \frac{1}{z_i!}\right) \exp\left\{\left(\sum_{i=1}^{B} z_i\right) \log x - Bx\right\},
\end{aligned}$$

we have that $h(z) = \prod_{i=1}^{B} \frac{1}{z_i!}, T(z) = \sum_{i=1}^{B} z_i, \theta_3 = B$. *Therefore, by Theorem A.2, when* $B = 1$,

$$\begin{aligned}
p(z \mid \theta_1, \theta_2, \theta_3) &= \frac{1}{z!} \frac{\frac{\beta^\alpha}{\Gamma(\alpha)}}{\frac{(\beta+1)^{(\alpha+z)}}{\Gamma(\alpha+z)}} \\
&= \frac{(\alpha + z - 1)!}{(\alpha - 1)! z!} \left(\frac{\beta}{\beta + 1}\right)^\alpha \left(\frac{1}{\beta + 1}\right)^z \\
&= \mathsf{NB}\left(z \mid \alpha, \frac{\beta}{\beta + 1}\right); \\
p(x \mid z, \theta_1, \theta_2, \theta_3) &= \frac{(\beta + 1)^{(\alpha+z)}}{\Gamma(\alpha + z)} \exp\{(\alpha + z - 1)\log(x) - (\beta + 1)x\} \\
&= \mathsf{Gam}(x \mid \alpha + z, \beta + 1).
\end{aligned}$$

However, when $B > 1$,

$$\begin{aligned}
p(z \mid \theta_1, \theta_2, \theta_3) &= \left(\prod_{i=1}^{B} \frac{1}{z_i!}\right) \frac{\frac{\beta^\alpha}{\Gamma(\alpha)}}{\frac{(\beta+B)^{(\alpha+\sum_{i=1}^{B} z_i)}}{\Gamma\left(\alpha+\sum_{i=1}^{B} z_i\right)}} \\
&\neq \prod_{i=1}^{B} \mathsf{NB}\left(z_i \mid \alpha, \frac{\beta}{\beta + B}\right); \\
p(x \mid z, \theta_1, \theta_2, \theta_3) &= \frac{(\beta + B)^{\left(\alpha+\sum_{i=1}^{B} z_i\right)}}{\Gamma\left(\alpha + \sum_{i=1}^{B} z_i\right)} \exp\left\{\left(\alpha - 1 + \sum_{i=1}^{B} z_i\right) \log(x) - (\beta + B)x\right\} \\
&= \mathsf{Gam}\left(x \mid \alpha + \sum_{i=1}^{B} z_i, \beta + B\right).
\end{aligned}$$

◁

Example A.4. *Suppose* $X \sim \mathsf{Gam}(\alpha, \beta)$. *Draw* $Z \sim \mathsf{Poiss}(\tau X)$, *where* $\tau \in (0, \infty)$ *is a tuning parameter. Let* $f(X) = Z$ *and* $g(X) = X$.
By writing

$$\mathsf{Gam}(x \mid \alpha, \beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \exp\{(\alpha - 1)\log x - \beta x\},$$

where $\mathsf{Gam}(\cdot \mid \alpha, \beta)$ *denotes the pdf of* $\mathsf{Gam}(\alpha, \beta)$*, we have that* $\theta_2 = \beta, \theta_1 = \alpha - 1, \phi(x) = \log x, A(\phi(x)) = \exp(\phi(x)) = \exp(\log x) = x, m(x) = 1$*. Therefore,* $H(\theta_1, \theta_2) = \frac{\theta_2^{(\theta_1+1)}}{\Gamma(\theta_1+1)}$*. Now since*

$$\begin{aligned}
\mathsf{Poiss}(z \mid \tau x) &= \frac{1}{z!} \exp\{z \log(\tau x) - \tau x\} \\
&= \frac{1}{z!} \exp\{z \log \tau + z \log x - \tau x\} \\
&= \frac{\tau^z}{z!} \exp\{z \log x - \tau x\},
\end{aligned}$$

we have that $h(z) = \frac{\tau^z}{z!}, T(z) = z, \theta_3 = \tau$*. Therefore, by Theorem A.2,*

$$\begin{aligned}
p(z \mid \theta_1, \theta_2, \theta_3) &= \frac{\tau^z}{z!} \frac{\frac{\beta^\alpha}{\Gamma(\alpha)}}{\frac{(\beta+\tau)^{(\alpha+z)}}{\Gamma(\alpha+z)}} \\
&= \frac{(\alpha+z-1)!}{(\alpha-1)!z!} \left(\frac{\beta}{\beta+\tau}\right)^\alpha \left(\frac{\tau}{\beta+\tau}\right)^z \\
&= \mathsf{NB}\left(z \mid \alpha, \frac{\beta}{\beta+\tau}\right); \\
p(x \mid z, \theta_1, \theta_2, \theta_3) &= \frac{(\beta+\tau)^{(\alpha+z)}}{\Gamma(\alpha+z)} \exp\left\{(\alpha+z-1)\log(x) - (\beta+\tau)x\right\} \\
&= \mathsf{Gam}\left(x \mid \alpha+z, \beta+\tau\right).
\end{aligned}$$

◁

Example A.5. *Assume for all* $i \in \{1, 2, \ldots, n\}$*,* $Y_i \overset{i.i.d.}{\sim} \mathsf{Gam}(\alpha, \exp(\beta^\top x_i))$*, where each* $x_i \in \mathbb{R}^d$ *is fixed. Following the data fission procedure in Example A.3, with* $B = 1$*, we have that for all* i*,* $f(Y_i) = Z_i, g(Y_i) = Y_i$*. In the selection phase of selective inference, for some fixed* $\lambda > 0$*, we can do model selection via the optimization below*

$$\begin{aligned}
\hat{\beta}_\lambda &= \underset{\beta \in \mathbb{R}^d}{\arg\min} - \sum_{i=1}^n \left(\log \mathsf{NB}\left(z_i \mid \alpha, \frac{\exp(\beta^\top x_i)}{\exp(\beta^\top x_i)+1}\right)\right) + \lambda \|\beta_1\| \\
&= \underset{\beta \in \mathbb{R}^d}{\arg\min} \left(\sum_{i=1}^n -\log \binom{z_i+\alpha-1}{z_i} - z_i \log\left(\frac{1}{\exp(\beta^\top x_i)+1}\right) - \alpha \log\left(\frac{\exp(\beta^\top x_i)}{\exp(\beta^\top x_i)+1}\right)\right) + \lambda \|\beta_1\|,
\end{aligned}$$

which is convex in β*. Denote the index set of nonzero entries in* $\hat{\beta}_\lambda$ *to be* $\mathcal{M}$ *and* $|\mathcal{M}| = d' \leq d$*.*
Using the selected features $\mathcal{M}$*, with* $x_{i,\mathcal{M}}$ *denoting the selected features for the* i*th observation, we can obtain the estimates* $\hat{\beta}_n(\mathcal{M})$ *via*

$$\begin{aligned}
&\hat{\beta}_n(\mathcal{M}) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} - \sum_{i=1}^n \left(\log \mathsf{Gam}(y_i \mid \alpha + z_i, \exp(\beta^\top x_{i,\mathcal{M}}) + 1)\right) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} \sum_{i=1}^n -(\alpha+z_i)\log\left(\exp(\beta^\top x_{i,\mathcal{M}})+1\right) + \log\Gamma(\alpha+z_i) - (\alpha+z_i-1)\log y_i + \left(\exp(\beta^\top x_{i,\mathcal{M}})+1\right) y_i,
\end{aligned}$$

which may be convex in β *but without an argmin, or non-convex (depending on the value of* α*). Therfore, we can instead use the working model*

$$\begin{aligned}
&\hat{\beta}_n(\mathcal{M}) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} - \sum_{i=1}^n \left(\log \mathsf{Gam}(y_i \mid \alpha + z_i, \exp(\beta^\top x_{i,\mathcal{M}}))\right) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} \sum_{i=1}^n -(\alpha+z_i)\log\left(\exp(\beta^\top x_{i,\mathcal{M}})\right) + \log\Gamma(\alpha+z_i) - (\alpha+z_i-1)\log y_i + \left(\exp(\beta^\top x_{i,\mathcal{M}})\right) y_i.
\end{aligned}$$

For this problem, we have that

$$\frac{\partial m_i}{\partial \eta_i} = \exp(\beta^\top x_{i,\mathcal{M}}), \quad m_i = \frac{\alpha + z_i}{\exp(\beta^\top x_{i,\mathcal{M}})}, \quad v_i = \frac{\alpha + z_i}{(\exp(\beta^\top x_{i,\mathcal{M}}))^2}.$$

Let D, V, M, G be diagonal matrices with

$$D_{ii} = \frac{\partial m_i}{\partial \eta_i}, \quad V_{ii} = v_i, \quad M_{ii} = m_i, \quad G_{ii} = g(y_i) - m_i,$$

then the plug-in estimator for variance becomes

$$\hat{H}_n^{-1} \hat{V}_n \hat{H}_n^{-1} = \left(X_{\mathcal{M}}^\top D^2 V^{-1} X_{\mathcal{M}}\right)^{-1} \left(X_{\mathcal{M}}^\top G^2 V^{-2} D^2 X_{\mathcal{M}}\right) \left(X_{\mathcal{M}}^\top D^2 V^{-1} X_{\mathcal{M}}\right)^{-1}.$$

◁

Example A.6. *Assume for all $i \in \{1, 2, \ldots, n\}$, $Y_i \overset{i.i.d.}{\sim} \mathsf{Gam}(\alpha, \exp(\beta^\top x_i))$, where each $x_i \in \mathbb{R}^d$ is fixed. Following the data fission procedure in Example A.4, we have that for all i, $f(Y_i) = Z_i, g(Y_i) = Y_i$. In the selection phase of selective inference, for some fixed $\lambda > 0$, we can do model selection via the optimization below*

$$\begin{aligned}
\hat{\beta}_\lambda &= \underset{\beta \in \mathbb{R}^d}{\arg\min} - \sum_{i=1}^n \left(\log \mathsf{NB}\left(z_i \mid \alpha, \frac{\exp(\beta^\top x_i)}{\exp(\beta^\top x_i) + \tau} \right) \right) + \lambda \|\beta_1\| \\
&= \underset{\beta \in \mathbb{R}^d}{\arg\min} \left(\sum_{i=1}^n - \log \binom{z_i + \alpha - 1}{z_i} - z_i \log \left(\frac{1}{\exp(\beta^\top x_i) + \tau} \right) - \alpha \log \left(\frac{\exp(\beta^\top x_i)}{\exp(\beta^\top x_i) + \tau} \right) \right) + \lambda \|\beta_1\|,
\end{aligned}$$

which is convex in β. Denote the index set of nonzero entries in $\hat{\beta}_\lambda$ to be $\mathcal{M}$ and $|\mathcal{M}| = d' \leq d$.
Using the selected features $\mathcal{M}$, with $x_{i,\mathcal{M}}$ denoting the selected features for the ith observation, we can obtain the estimates $\hat{\beta}_n(\mathcal{M})$ via

$$\begin{aligned}
&\hat{\beta}_n(\mathcal{M}) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} - \sum_{i=1}^n \left(\log \mathsf{Gam}(y_i \mid \alpha + z_i, \exp(\beta^\top x_{i,\mathcal{M}}) + \tau) \right) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} \sum_{i=1}^n -(\alpha + z_i) \log \left(\exp(\beta^\top x_{i,\mathcal{M}}) + \tau \right) + \log \Gamma(\alpha + z_i) - (\alpha + z_i - 1) \log y_i + \left(\exp(\beta^\top x_{i,\mathcal{M}}) + \tau \right) y_i,
\end{aligned}$$

which may be convex in β but without an argmin, or non-convex (depending on the value of α). Therefore, we can instead use the working model

$$\begin{aligned}
&\hat{\beta}_n(\mathcal{M}) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} - \sum_{i=1}^n \left(\log \mathsf{Gam}(y_i \mid \alpha + z_i, \exp(\beta^\top x_{i,\mathcal{M}})) \right) \\
&= \underset{\beta \in \mathbb{R}^{d'}}{\arg\min} \sum_{i=1}^n -(\alpha + z_i) \log \left(\exp(\beta^\top x_{i,\mathcal{M}}) \right) + \log \Gamma(\alpha + z_i) - (\alpha + z_i - 1) \log y_i + \left(\exp(\beta^\top x_{i,\mathcal{M}}) \right) y_i.
\end{aligned}$$

For this problem, we have that

$$\frac{\partial m_i}{\partial \eta_i} = \exp(\beta^\top x_{i,\mathcal{M}}), \quad m_i = \frac{\alpha + z_i}{\exp(\beta^\top x_{i,\mathcal{M}})}, \quad v_i = \frac{\alpha + z_i}{(\exp(\beta^\top x_{i,\mathcal{M}}))^2}.$$

Let D, V, M, G be diagonal matrices with

$$D_{ii} = \frac{\partial m_i}{\partial \eta_i}, \quad V_{ii} = v_i, \quad M_{ii} = m_i, \quad G_{ii} = g(y_i) - m_i,$$

then the plug-in estimator for variance becomes

$$\hat{H}_n^{-1}\hat{V}_n\hat{H}_n^{-1} = \left(X_{\mathcal{M}}^\top D^2 V^{-1} X_{\mathcal{M}}\right)^{-1}\left(X_{\mathcal{M}}^\top G^2 V^{-2} D^2 X_{\mathcal{M}}\right)\left(X_{\mathcal{M}}^\top D^2 V^{-1} X_{\mathcal{M}}\right)^{-1}.$$

◁

No need to ensure $\exp(\beta^\top x_{i,\mathcal{M}})$ stays above 1 or τ, since mean already have same support?